

Introduction to Computation Technologies in Deep Learning

Kai Jia
jiakai@megvii.com

Megvii Inc.

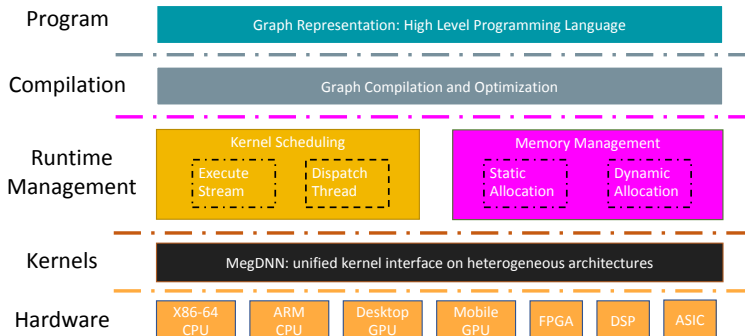
May. 30, 2018



- 1 Symbolic Computation
 - Representation
 - Execution & Optimization
- 2 Dense Numerical Computation
 - CPU Computation
 - Other Computation Devices
 - Computation & Memory Gap
- 3 Distributed Computation
 - System
 - Optimzation Algorithms
 - Communication Algorithms

Overview of a Deep Learning Framework


MegBrain Architecture



Computation Graph

$$y = \text{conv}(x, w) + b$$

Computation Graph



Graph Structure

Variable

- Corresponding to a tensor with concrete numerical values during graph execution

Graph Structure

Variable

- Corresponding to a tensor with concrete numerical values during graph execution
- Carrying some *attributes*:

DType Data type, like `int8` and `float32`.

Graph Structure

Variable

- Corresponding to a tensor with concrete numerical values during graph execution
- Carrying some *attributes*:

DType Data type, like `int8` and `float32`.

Shape Like `(5, 3)` for FC weight, `(128, 512, 7, 7)` for feature maps.

Example

Variable shape inference facilitates automatic weight initialization:

```
assert x.shape == (128, 50)
y = fully_connected(x, output_dim=100)
assert y.shape == (128, 100)
```

The weight matrix of this `FullyConnected` operator can be initialized to `np.random.normal((50, 100))`.

Graph Structure


Operators & Edges

Operator

- Operators & variables form a bipartite graph
- Operators define the computation to be applied on input variables

Edge

- Data dependency: read input data
- Control dependency: require input operator to have finished



Operator Granularity

Trade-off between flexibility and ease-to-optimize

Category	Example	Advantage	Framework
Coarse-grained	$y = \text{BatchNorm}(x)$	Parsimony; Easy performance tuning	Caffe
Fine-grained	$y = \frac{x - \text{mean}(x)}{\text{std}(x)}$	Flexibility	Theano

Operator Granularity

Trade-off between flexibility and ease-to-optimize

Category	Example	Advantage	Framework
Coarse-grained	$y = \text{BatchNorm}(x)$	Parsimony; Easy performance tuning	Caffe
Fine-grained	$y = \frac{x - \text{mean}(x)}{\text{std}(x)}$	Flexibility	Theano

Our philosophy:

- Prefer **flexibility**: this can not be changed once the framework has been designed
- Utilize **multi-level API** for simplifying graph representation
- Speed can be continuously improved by **graph optimizer**

Auto Differentiation

Calculus 101: The Chain Rule

Straight-forward approach: each operator provides two methods: fprop and bprop.

```
inline Dtype Forward(const vector<Blob<Dtype>*>& bottom,  
    const vector<Blob<Dtype>*>& top);
```

```
inline void Backward(const vector<Blob<Dtype>*>& top,  
    const vector<bool>& propagate_down,  
    const vector<Blob<Dtype>*>& bottom);
```

Auto Differentiation

Calculus 101: The Chain Rule

Straight-forward approach: each operator provides two methods: fprop and bprop.

Limitations

- Graph optimizer can not be uniformly applied on both forward and backward passes
- Difficult to implement gradient of gradient ($\frac{\partial f(\frac{\partial L}{\partial x})}{\partial y}$, in WGAN training ¹) or higher-order gradients ($\frac{\partial^2 L}{\partial x^2}$).
- Difficult to modify/manipulate gradients (e.g. for low-bit training).

¹Ishaan Gulrajani et al. "Improved training of wasserstein gans". In: *arXiv preprint arXiv:1704.00028* (2017).

Auto Differentiation

Calculus 101: The Chain Rule

Unified approach: extending the graph with operators computing gradients of specific variables, via the chain rule.

$$y_1, \dots, y_m = f(x_1, \dots, x_n)$$

Gradient operator g for f .

Auto Differentiation

Calculus 101: The Chain Rule

Unified approach: extending the graph with operators computing gradients of specific variables, via the chain rule.


$$y_1, \dots, y_m = f(x_1, \dots, x_n)$$

Gradient operator g for f :

$$\frac{\partial L}{\partial x_1}, \dots, \frac{\partial L}{\partial x_n} = g\left(\frac{\partial L}{\partial y_1}, \dots, \frac{\partial L}{\partial y_m}, x_1, \dots, x_n, y_1, \dots, y_m\right)$$

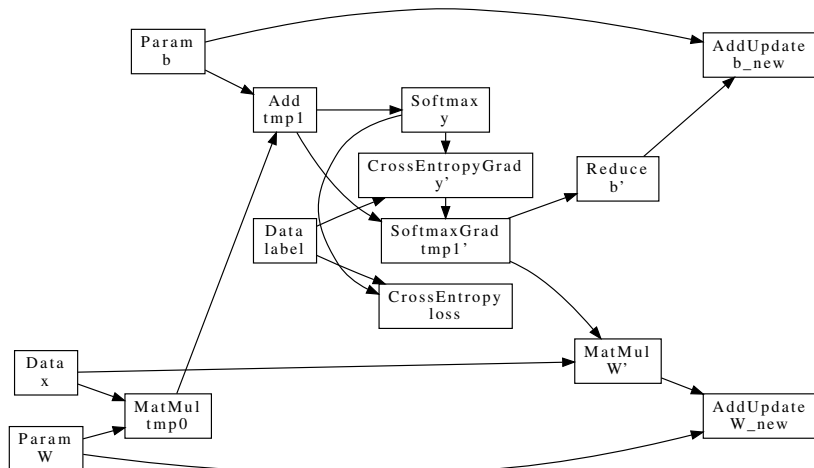
Auto Differentiation

Calculus 101: The Chain Rule



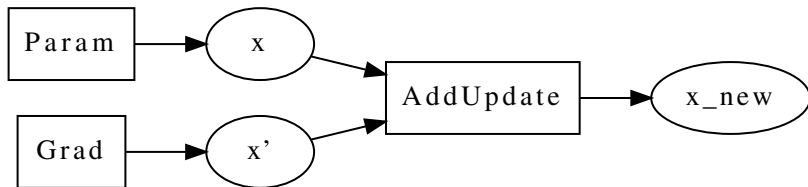
Auto Differentiation

Calculus 101: The Chain Rule



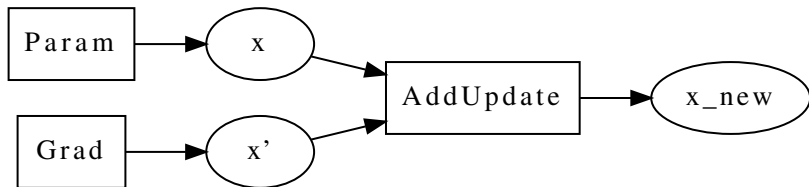
Mutable State

Enabling param updates to be expressed in graphs



Mutable State

Enabling param updates to be expressed in graphs




Note

- x and x_new share the underlying storage and should not be simultaneously read by one operator. Equivalently speaking, `AddUpdate` separates the graph.

Mutable State

Enabling param updates to be expressed in graphs



Note

- x and x_{new} share the underlying storage and should not be simultaneously read by one operator. Equivalently speaking, `AddUpdate` separates the graph.
- Readers of x must have finished (impl. by control dependency)

Symbolic Shape

Enabling computation involving tensor shapes

Example

Maxout¹:




Image from ²

²Hai Dai Nguyen, Anh Duc Le, and Masaki Nakagawa. "Recognition of Online Handwritten Math Symbols Using Deep Neural Networks". In: *IEICE Trans. Inf. & Syst.* 99.12 (2016), pp. 3110–3118.

¹Ian J Goodfellow et al. "Maxout networks". In: *arXiv preprint*

Symbolic Shape

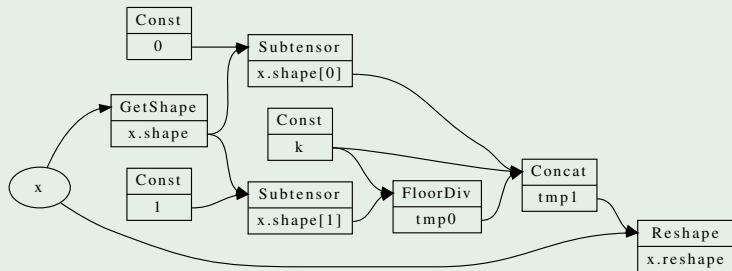
Enabling computation involving tensor shapes

Example

Maxout¹:

```
y = x.reshape(x.shape[0], x.shape[1] // k, k).max(axis=2)
```

where `x.shape` is also a symbol whose value is evaluated at runtime, so the computation adapts to different input shapes.



¹ Ian J Goodfellow et al. "Maxout networks". In: *arXiv preprint*

Symbolic Shape

Enabling computation involving tensor shapes

Example

Maxout¹:

```
y = x.reshape(x.shape[0], x.shape[1] // k, k).max(axis=2)
```

- Helps dealing with non-constant batch size or input image size
- Requires dynamic shape support: some shapes may remain unknown until graph execution

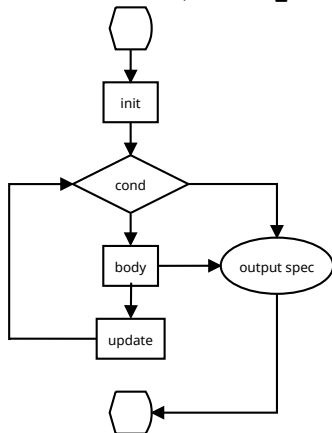
¹Ian J Goodfellow et al. "Maxout networks". In: *arXiv preprint arXiv:1302.4389* (2013).

Control Flow Operators

Towards universal computation (in theory)

Loop operator:

`scan` in Theano, `while_loop` in TensorFlow and `loop` in MegBrain.




Control Flow Operators

Towards universal computation (in theory)

Loop operator:

`scan` in Theano, `while_loop` in TensorFlow and `loop` in MegBrain.



- With symbolic shapes and control flow operators, a computation graph is **Turing-complete**!
- Useful for RNN and iterative algorithms

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Static Computation Graph

- **Unfamiliar programming model:**

Stateless, functional: `y = x.setsub[1:3](xs)`

rather than imperative: `x[1:3] = xs`

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Static Computation Graph

- **Unfamiliar programming model:**

Stateless, functional: `y = x.setsub[1:3](xs)`

rather than imperative: `x[1:3] = xs`

- **Difficult to debug:** code is written for graph construction but tensor values can only be known during graph execution

`y = printop(y)` rather than `print(y)`

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

- Auto differentiation: keep symbolic track of computation path

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

- Auto differentiation: keep symbolic track of computation path
- Drawbacks:
 - hard to optimize: lack of global information
 - hard to deploy: graph depends on code

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Graph Execution

Separation of representation and execution allows abstraction of hardware details

- Map from variables to tensor values

Graph Execution


Separation of representation and execution allows abstraction of hardware details

- Map from variables to tensor values
- Map from operators to kernels on some specific architecture

Graph Execution

Separation of representation and execution allows abstraction of hardware details

- Map from variables to tensor values
- Map from operators to kernels on some specific architecture
- Kernels are scheduled according to topological order



(1, 2, 3, 4) or (1, 3, 2, 4)

Optimizing by Graph Transformation

- Expression simplifying: $x + 1 - 2 + x \Rightarrow 2x - 1$

Optimizing by Graph Transformation

- Expression simplifying: $x + 1 - 2 + x \Rightarrow 2x - 1$
- Operation reordering according to shapes
tensors: x and y
scalars: a and b
 $x + a + y + b \Rightarrow a + b + x + y$

Optimizing by Graph Transformation

- Expression simplifying: $x + 1 - 2 + x \Rightarrow 2x - 1$
- Operation reordering according to shapes
tensors: x and y
scalars: a and b
 $x + a + y + b \Rightarrow a + b + x + y$
- Operator fusion: $x \cdot y + z \Rightarrow \text{fma}(x, y, z)$
 - static fusion: predefined fusion rules
 - dynamic fusion: Just-in-time compilation (JIT) for actual computation graph

Runtime Memory Management

- Baseline: reference counting + some classical memory allocator

Runtime Memory Management


- Baseline: reference counting + some classical memory allocator
 - Readonly forwarding: reuse input storage for operators like reshape and subtensor
 - Writable forwarding (a.k.a. inplace operation): overwrite input storage
- Caution:** must ensure no other readers exist (i.e. refcnt equals 1)

Sublinear Memory

Trade time for memory

Observation

- Long-term dependency for gradient computing consumes lots of memory.
- Assume $x_{i+1} = \text{conv}(x_i, w_i)$, then x_{i+1} can only be discarded after $\frac{\partial L}{\partial w_i}$ and $\frac{\partial L}{\partial x_i}$ have been computed.




Sublinear Memory

Trade time for memory

Observation

- Long-term dependency for gradient computing consumes lots of memory.
- Assume $x_{i+1} = \text{conv}(x_i, w_i)$, then x_{i+1} can only be discarded after $\frac{\partial L}{\partial w_i}$ and $\frac{\partial L}{\partial x_i}$ have been computed.



Sublinear Memory

Trade time for memory

Observation

- Long-term dependency for gradient computing consumes lots of memory.
- Assume $x_{i+1} = \text{conv}(x_i, w_i)$, then x_{i+1} can only be discarded after $\frac{\partial L}{\partial w_i}$ and $\frac{\partial L}{\partial x_i}$ have been computed.

Method

- Split the sequence into blocks consisting of consecutive operators and keep only the first variable in any block
- Recompute internal values in a block when gradient is needed
- In the above example, discard x_{km+j} for all $0 < j < m$ and recompute them when needed.

Sublinear Memory

Trade time for memory

Reduce memory usage to $O(\sqrt{n})$ with extra $O(n)$ time cost in the ideal case.

For a graph with 10000 convolutions and their gradients:

comp_node	alloc	lower_bound	upper_bound
gpu0:0	15624.37MiB(16383336448bytes)	15624.37MiB(100.00%)	31889.13MiB(204.10%)
comp_node	alloc	lower_bound	upper_bound
gpu0:0	173.03MiB(181430784bytes)	168.76MiB(97.53%)	47251.78MiB(27309.08%)

Note: this idea is also published in².

²Tianqi Chen et al. "Training deep nets with sublinear memory cost". In: *arXiv preprint arXiv:1604.06174* (2016).

- 1 Symbolic Computation
 - Representation
 - Execution & Optimization
- 2 Dense Numerical Computation
 - CPU Computation
 - Other Computation Devices
 - Computation & Memory Gap
- 3 Distributed Computation
 - System
 - Optimzation Algorithms
 - Communication Algorithms

Instruction: The Hardware/Software Interface

What is a program?

```
int sum(int *x) {  
    return x[0] + x[1];  
}
```

000000000000000000 <sum>:

0: 8b 47 04

3: 03 07

5: c3

mov 0x4(%rdi),%eax

add (%rdi),%eax

retq

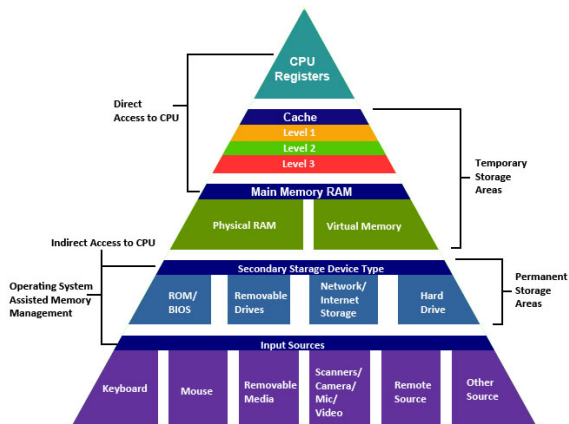
Modern CPU Technologies

Instr. No.	Pipeline Stage						
	IF	ID	EX	MEM	WB		
1							
2							
3							
4							
5							
Clock Cycle	1	2	3	4	5	6	7


- **Pipeline** ³
- **Superpipelining** increases stage number and simplifies each stage
- **Superscalar** dispatches multiple instructions to implement instruction-level parallelism
- **Out-of-order execution** executes according to availability of input data rather than original program order

³ Image from https://en.wikipedia.org/wiki/Instruction_pipelining

Memory Hierarchy



Static Random-Access Memory (SRAM)




- ① Fast
- ② Low power consumption
- ③ No refresh circuit

Image from https://en.wikipedia.org/wiki/Static_random-access_memory

RAM Implementation

Dynamic Random-Access Memory (DRAM)



Advantages:

- ① High density
- ② Cheap


Refresh: periodically read blocks and write back.

Image from <http://docencia.ac.upc.edu/master/MIRI/NCD/docs/04-Memory%20Structures-2.pdf>

Cache Hierarchy

A hierarchical design for better trade-off between memory capacity and latency.

eDRAM Based Cache



26

Intel Next Generation Microarchitecture Code Name Skylake

IDF15
INTEL DEVELOPER FORUM

Image from

<https://www.anandtech.com/show/9582/intel-skylake-mobile-desktop-launch-architecture-analysis/5>

Cache Hierarchy

A hierarchical design for better trade-off between memory capacity and latency.

Core i7 Xeon 5500 Series

L1 hit	~ 4 cycles
L2 hit	~ 10 cycles
L3 hit line unshared	~ 40 cycles
L3 hit, shared line in another core	~ 65 cycles
L3 hit, modified in another core	~ 75 cycles
Remote L3	~ 100 – 300 cycles
Local DRAM	~ 60 ns
Remote DRAM	~ 100 ns

source:

https://software.intel.com/sites/products/collateral/hpc/vtune/performance_analysis_guide.pdf

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

- Associativity

Number of different tags to be kept under the same index

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

- Associativity

Number of different tags to be kept under the same index

- VIPT Addressing

Virtually indexed, physically tagged (VIPT)

- Solve aliasing problem
- Simultaneous cache and TLB lookup

Interesting reading: <http://igoro.com/archive/gallery-of-processor-cache-effects/>

CPU Cache Structure

Example

```
$ grep -m1 name /proc/cpuinfo
model name      : Intel(R) Core(TM) i5-6200U CPU @ 2.30GHz

$ cat /sys/devices/system/cpu/cpu0/cache/index0/{size,ways_of_associativity,coherency_line_size}
32K
8
64
```

- block offset: $\log_2 64 = 6\text{bit}$
- index: $\log_2(32\text{KiB}/64\text{B}/8) = 6\text{bit}$
- tag: $52 - 6 - 6 = 40\text{bit}$ (48-bit virtual memory and 52-bit physical memory)

Note that *block offset* and *index* together take 12 bits, which is equal to page size (4KiB), so VIPT can be easily implemented.

SIMD

Single instruction, multiple data

Store multiple data items in one register and process them in a single instruction.

Calculation of theoretical FLOPS⁴

$$FLOPS = f \cdot w \cdot IPC$$

f : frequency

w : SIMD width (number of floats per register)

IPC : SIMD instructions per cycle

⁴floating point operations per second

SIMD

Single instruction, multiple data

Example

Intel® CPUs usually have $IPC = 2$. However if FMA is supported, IPC should be counted as 4 since 2 FMA instructions is essentially 4 floating point operations.

Intel® Xeon® Platinum 8180M⁴

# of Cores	28
Processor Base Frequency	2.50 GHz
Max Turbo Frequency	3.80 GHz
# of AVX-512 FMA Units	2

$$\begin{aligned}
 FLOPS &= 3.8Gcyc/s \times 4instr/cyc \times 16float/instr \\
 &= 243.2GFLOPS
 \end{aligned}$$

$$FLOPS_{TOT} = FLOPS \times 28 = 6.8TFLOPS$$

⁴ data available at

A MatMul Example

```
void matmul(float *a, float *b, float *c, int n) {  
    for (int i = 0; i < n; ++ i) {  
        for (int j = 0; j < n; ++ j) {  
            float sum = 0;  
            for (int k = 0; k < n; ++ k) {  
                sum += a[i * n + k] * b[k * n + j];  
            }  
            c[i * n + j] = sum;  
        }  
    }  
}
```

A MatMul Example

```
void matmul(float *a, float *b, float *c, int n) {  
    for (int i = 0; i < n; ++ i) {  
        for (int j = 0; j < n; ++ j) {  
            float sum = 0;  
            for (int k = 0; k < n; ++ k) {  
                sum += a[i * n + k] * b[k * n + j];  
            }  
            c[i * n + j] = sum;  
        }  
    }  
}
```

Swap the loops on j and k

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

NVIDIA GPU

A single instruction, multiple thread architecture




Image from⁵

⁵Marco Nobile et al. "cuTauLeaping: A GPU-Powered Tau-Leaping Stochastic Simulator for Massive Parallel Analyses of Biological Systems". In: 9 (Mar. 2014), e91963.

NVIDIA GPU

A single instruction, multiple thread architecture

```
__global__ void add(float *a, float *b, float *c, int n) {  
    int id = blockIdx.x*blockDim.x+threadIdx.x;  
    if (id < n)  
        c[id] = a[id] + b[id];  
}
```

- **Memory Coalescing**



Adjacent threads access adjacent memories simultaneously

- **Parallelism**

Divide the total work among many tiny threads

NVIDIA GPU

A single instruction, multiple thread architecture



Tesla V100 for NVLink

- 15.7 *TFLOPS* for single-precision
- 125 *TFLOPS* for half-precision

Image from <https://arstechnica.com/gadgets/2017/05/nvidia-tesla-v100-gpu-details/>

The Trend

- On cloud: high density computation
e.g. Google TPU 3.0 pods are claimed to achieve 100PFLOPS
- On edge: low precision
e.g. int8 in cDSP supported by Qualcomm's SNPE
- Automatic kernel tuning & generation
An active research area. Typical projects include Halide⁵, TVM⁶ and TensorComprehension⁷

⁵<http://halide-lang.org/>

⁶<http://tvm-lang.org/>

⁷<https://facebookresearch.github.io/TensorComprehensions/>

1 Symbolic Computation

- Representation
- Execution & Optimization


2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Unbalanced Development of Processor and Memory



Graph from⁸

⁸Carlos Carvalho. "The gap between processor and memory speeds". In: *Proc. of IEEE International Conference on Control and Automation*. 2002.

Challenges from NN Architecture

Computation-sparse structure seems to be beneficial.

Architecture	Computation	Memory
Small kernel	$\frac{k2^2}{k1^2}$	Param $\frac{k1^2}{k2^2}$
Large stride	$\frac{1}{s^2}$	Output $\frac{1}{s^2}$
Group/depthwise conv	$\frac{1}{g^2}$	Param $\frac{1}{g}$
Shuffle/concat	0	1

Roofline Model

A visualization method to characterize computation/memory

Performance P (FLOPS) is approximately a function of arithmetic intensity I (FLOP/byte) for a particular architecture⁹.

Naïve Roofline

$$P = \min \left\{ \begin{array}{l} \pi \\ \beta \times I \end{array} \right.$$

where π is the peak computing performance and β is the peak bandwidth.

⁹Samuel Webb Williams. *Auto-tuning performance on multicore computers*. University of California, Berkeley, 2008.

Roofline Model

A visualization method to characterize computation/memory




Image from https://en.wikipedia.org/wiki/Roofline_model

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation


- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Communication System: Single Node

- PCI-e: connection between GPUs, network adaptors and others
 - Switches may be needed
 - 985 MiB/s each PCI-e 3.0 lane
 - LGA-2011 socket: 40 lanes
- NVLink: GPU interconnect by NVIDIA



Communication System: LAN

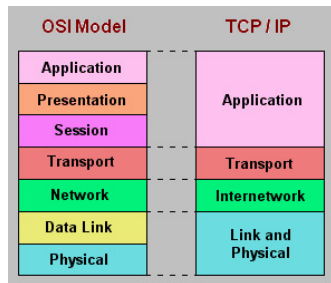


Image from <http://www.just2good.co.uk/tcpipStack.php>

- Ethernet: 10M to 100G, latency 100 - 20 μ s
- InfiniBand: 2.5 to 250G, latency 5 - 0.5 μ s

RDMA

Remote Direct Memory Access

Bypass the TCP/IP stack and free CPU from handling packets.

RDMA

Remote Direct Memory Access

Bypass the TCP/IP stack and free CPU from handling packets.

- RoCE: RDMA over Converged Ethernet
- InfiniBand: RDMA supported
- NVIDIA GPUDirect: RDMA between GPUs

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized


$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$

Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$




Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$




Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$

SGD step 2



Asynchronous SGD

Basics

Each worker has an outdated local copy of params and updates central param storage asynchronously. Friendly for parallel speedup.




Image from⁹

⁹Jeffrey Dean et al. "Large scale distributed deep networks". In: *NIPS*. 2012, pp. 1223–1231.

Asynchronous SGD

Difficulties

ASGD is not equivalent to SGD and it is hard to tune due to noisy gradients. Many works exist on analyzing convergence and improving performance¹⁰¹¹¹².




Image from¹³

¹⁰Xiangru Lian et al. "Asynchronous Parallel Stochastic Gradient for Nonconvex Optimization". In: *NIPS*. 2015, pp. 2737–2745.

¹¹Wei Zhang et al. "Staleness-aware async-SGD for Distributed Deep Learning". In: *IJCAI*. 2016, pp. 2350–2356.

¹²Sixin Zhang, Anna E Choromanska, and Yann LeCun. "Deep learning with elastic averaging SGD". In: *NIPS*. 2015, pp. 685–693.

¹³Jianmin Chen et al. "Revisiting distributed synchronous SGD". In: *arXiv preprint arXiv:1604.00981* (2016).

Synchronous SGD

Improvements

- Reduce communication by compressing gradients¹⁴

¹⁴Ryota Tomioka and Milan Vojnovic. "QSGD: Communication-Efficient Stochastic Gradient Descent, with Applications to Training Neural Networks". In: *arXiv preprint arXiv:1610.02132* (2016).

Synchronous SGD

Improvements

- Reduce communication by compressing gradients¹⁴

$$Q_s(v_i) = \|\mathbf{v}\|_2 \cdot \text{sgn}(v_i) \cdot \xi_i(\mathbf{v}, s)$$

¹⁴Ryota Tomioka and Milan Vojnovic. "QSGD: Communication-Efficient Stochastic Gradient Descent, with Applications to Training Neural Networks". In: *arXiv preprint arXiv:1610.02132* (2016).

Synchronous SGD

Improvements

- Reduce communication by compressing gradients¹⁴
- Handle straggling workers by backup workers¹⁵
Use $N + b$ workers but only receive gradients from any N of them and do not wait for the slowest b workers.

¹⁴[Ryota Tomioka and Milan Vojnovic](#). “QSGD: Communication-Efficient Stochastic Gradient Descent, with Applications to Training Neural Networks”. In: *arXiv preprint arXiv:1610.02132* (2016).

¹⁵[Jianmin Chen et al.](#) “Revisiting distributed synchronous SGD”. In: *arXiv preprint arXiv:1604.00981* (2016).

Synchronous SGD

Improvements

- Reduce communication by compressing gradients¹⁴
- Handle straggling workers by backup workers¹⁵
- Ensure performance by careful hyperparam tuning¹⁶
8192 minibatch size on 256 GPUs:

$$\hat{\eta} = k\eta$$
$$m = \frac{\eta_{t+1}}{\eta_t}$$

¹⁴Ryota Tomioka and Milan Vojnovic. “QSGD: Communication-Efficient Stochastic Gradient Descent, with Applications to Training Neural Networks”. In: *arXiv preprint arXiv:1610.02132* (2016).

¹⁵Jianmin Chen et al. “Revisiting distributed synchronous SGD”. In: *arXiv preprint arXiv:1604.00981* (2016).

¹⁶Priya Goyal et al. “Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour”. In: *arXiv preprint arXiv:1706.02677* (2017).

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap


3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms


MPI Primitives

Collective communication routines in MPI are common in distributed DL


MPI_Bcast




MPI_Scatter




MPI_Gather




MPI_Allgather



MPI_Reduce



MPI_Allreduce



An AllReduce Algorithm

Assume message size K and number of workers N

- Reduce to a worker (assume W_{N-1} here):
 W_i sends to W_{i+1} at step i ; communication per worker is N
- Broadcast from a worker: as above

An AllReduce Algorithm


Assume message size K and number of workers N

- Reduce to a worker (assume W_{N-1} here):
 W_i sends to W_{i+1} at step i ; communication per worker is N
- Broadcast from a worker: as above
- AllReduce:
 - 1 Split the message into N parts
 - 2 Reduce the i th part to W_i ; all reductions run in parallel
 - 3 Broadcast each reduced part to all workers in parallel

Communication cost for each worker is $2(N-1)\frac{K}{N}$,
independent of N .

More details and discussions are given in^{17 18}.

¹⁷Rajeev Thakur, Rolf Rabenseifner, and William Gropp. "Optimization of collective communication operations in MPICH". In: *The International Journal of High Performance Computing Applications* 19.1 (2005), pp. 49–66.

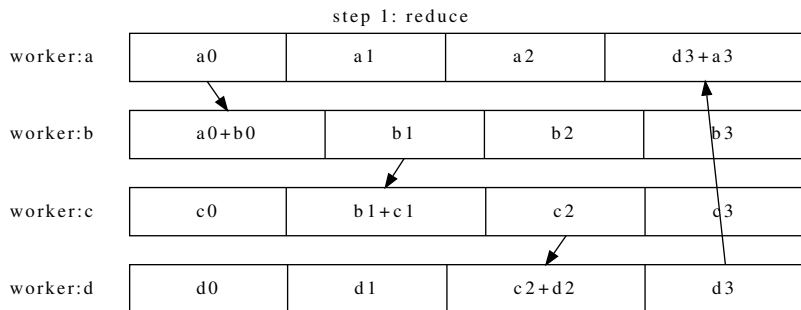
¹⁸<http://research.baidu.com/bringing-hpc-techniques-deep-learning> 

An AllReduce Algorithm


step 0: init

worker:a	a0	a1	a2	a3
worker:b	b0	b1	b2	b3
worker:c	c0	c1	c2	c3
worker:d	d0	d1	d2	d3


An AllReduce Algorithm



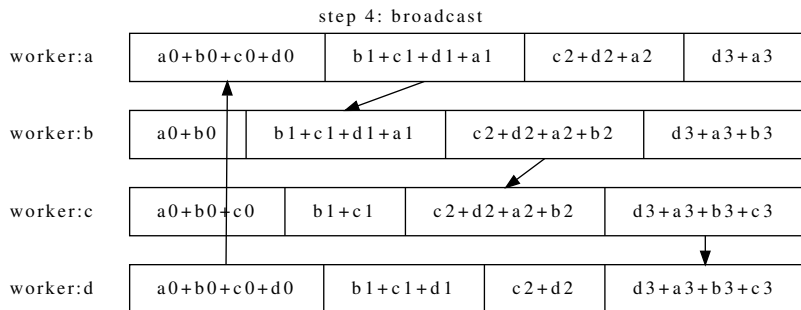
An AllReduce Algorithm




An AllReduce Algorithm



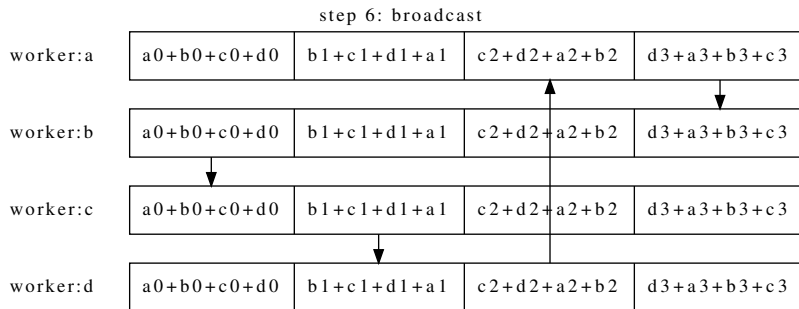
An AllReduce Algorithm



An AllReduce Algorithm



An AllReduce Algorithm



Thanks!

Questions and feedback are welcome :)