

Introduction to Computation Technologies in Deep Learning

Kai Jia
jiakai@megvii.com

Megvii Inc.

May. 30, 2018



1 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

Instruction: The Hardware/Software Interface

What is a program?

```
int sum(int *x) {  
    return x[0] + x[1];  
}
```

000000000000000000 <sum>:

0: 8b 47 04

3: 03 07

5: c3

mov 0x4(%rdi),%eax

add (%rdi),%eax

retq

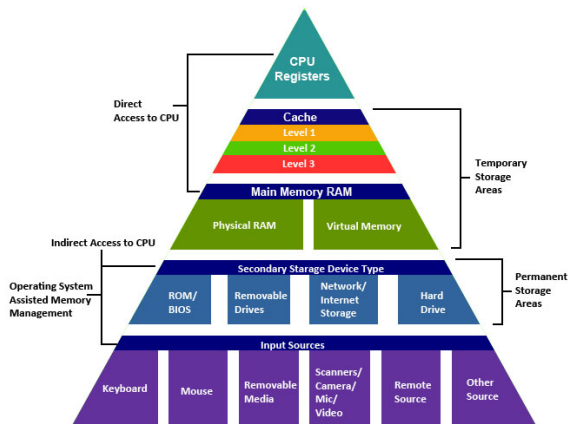
Modern CPU Technologies

Instr. No.	Pipeline Stage						
	IF	ID	EX	MEM	WB		
1							
2							
3							
4							
5							
Clock Cycle	1	2	3	4	5	6	7

- **Pipeline**¹
- **Superpipelining** increases stage number and simplifies each stage
- **Superscalar** dispatches multiple instructions to implement instruction-level parallelism
- **Out-of-order execution** executes according to availability of input data rather than original program order


¹ Image from https://en.wikipedia.org/wiki/Instruction_pipelining

Memory Hierarchy



RAM Implementation

Static Random-Access Memory (SRAM)



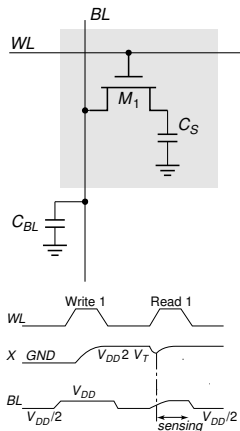
Advantages:

- 1 Fast
- 2 Low power consumption
- 3 No refresh circuit

Image from https://en.wikipedia.org/wiki/Static_random-access_memory

RAM Implementation

Dynamic Random-Access Memory (DRAM)



Advantages:

- ① High density
- ② Cheap

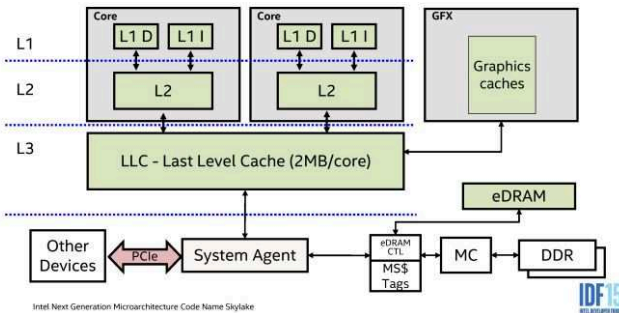
Refresh: periodically read blocks and write back.

Image from <http://docencia.ac.upc.edu/master/MIRI/NCD/docs/04-Memory%20Structures-2.pdf>

Cache Hierarchy

A hierarchical design for better trade-off between memory capacity and latency.

eDRAM Based Cache



26

Intel Next Generation Microarchitecture Code Name Skylake

IDF15
INTEL DEVELOPER FORUM

Image from

<https://www.anandtech.com/show/9582/intel-skylake-mobile-desktop-launch-architecture-analysis/5>

Cache Hierarchy

A hierarchical design for better trade-off between memory capacity and latency.

Core i7 Xeon 5500 Series

L1 hit	~ 4 cycles
L2 hit	~ 10 cycles
L3 hit line unshared	~ 40 cycles
L3 hit, shared line in another core	~ 65 cycles
L3 hit, modified in another core	~ 75 cycles
Remote L3	~ 100 – 300 cycles
Local DRAM	~ 60 ns
Remote DRAM	~ 100 ns

source:

https://software.intel.com/sites/products/collateral/hpc/vtune/performance_analysis_guide.pdf

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

- Associativity

Number of different tags to be kept under the same index

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

- Associativity

Number of different tags to be kept under the same index

- VIPT Addressing

Virtually indexed, physically tagged (VIPT)

- Solve aliasing problem
- Simultaneous cache and TLB lookup

Interesting reading: <http://igoro.com/archive/gallery-of-processor-cache-effects/>

CPU Cache Structure

Example

```
$ grep -m1 name /proc/cpuinfo
model name      : Intel(R) Core(TM) i5-6200U CPU @ 2.30GHz

$ cat /sys/devices/system/cpu/cpu0/cache/index0/{size,ways_of_associativity,coherency_line_size}
32K
8
64
```

- block offset: $\log_2 64 = 6\text{bit}$
- index: $\log_2(32\text{KiB}/64\text{B}/8) = 6\text{bit}$
- tag: $52 - 6 - 6 = 40\text{bit}$ (48-bit virtual memory and 52-bit physical memory)

Note that *block offset* and *index* together take 12 bits, which is equal to page size (4KiB), so VIPT can be easily implemented.

SIMD

Single instruction, multiple data

Store multiple data items in one register and process them in a single instruction.

Calculation of theoretical FLOPS²

$$FLOPS = f \cdot w \cdot IPC$$

f : frequency

w : SIMD width (number of floats per register)

IPC : SIMD instructions per cycle

²floating point operations per second

SIMD

Single instruction, multiple data

Example

Intel® CPUs usually have $IPC = 2$. However if FMA is supported, IPC should be counted as 4 since 2 FMA instructions is essentially 4 floating point operations.

Intel® Xeon® Platinum 8180M²

# of Cores	28
Processor Base Frequency	2.50 GHz
Max Turbo Frequency	3.80 GHz
# of AVX-512 FMA Units	2

$$\begin{aligned}
 FLOPS &= 3.8 \text{Gcyc/s} \times 4 \text{instr/cyc} \times 16 \text{float/instr} \\
 &= 243.2 \text{GFLOPS}
 \end{aligned}$$

$$FLOPS_{TOT} = FLOPS \times 28 = 6.8 \text{TFLOPS}$$

² data available at

A MatMul Example

```
void matmul(float *a, float *b, float *c, int n) {  
    for (int i = 0; i < n; ++ i) {  
        for (int j = 0; j < n; ++ j) {  
            float sum = 0;  
            for (int k = 0; k < n; ++ k) {  
                sum += a[i * n + k] * b[k * n + j];  
            }  
            c[i * n + j] = sum;  
        }  
    }  
}
```


A MatMul Example

```
void matmul(float *a, float *b, float *c, int n) {  
    for (int i = 0; i < n; ++ i) {  
        for (int j = 0; j < n; ++ j) {  
            float sum = 0;  
            for (int k = 0; k < n; ++ k) {  
                sum += a[i * n + k] * b[k * n + j];  
            }  
            c[i * n + j] = sum;  
        }  
    }  
}
```

Swap the loops on j and k

1 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

NVIDIA GPU

A single instruction, multiple thread architecture

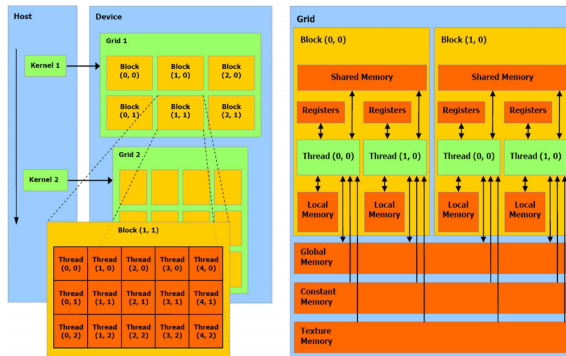


Image from³

³Marco Nobile et al. "cuTauLeaping: A GPU-Powered Tau-Leaping Stochastic Simulator for Massive Parallel Analyses of Biological Systems". In: 9 (Mar. 2014), e91963.

NVIDIA GPU

A single instruction, multiple thread architecture

```
__global__ void add(float *a, float *b, float *c, int n) {  
    int id = blockIdx.x*blockDim.x+threadIdx.x;  
    if (id < n)  
        c[id] = a[id] + b[id];  
}
```

- **Memory Coalescing**



Adjacent threads access adjacent memories simultaneously

- **Parallelism**

Divide the total work among many tiny threads

NVIDIA GPU

A single instruction, multiple thread architecture



Tesla V100 for NVLink

- 15.7 *TFLOPS* for single-precision
- 125 *TFLOPS* for half-precision

Image from <https://arstechnica.com/gadgets/2017/05/nvidia-tesla-v100-gpu-details/>

The Trend

- On cloud: high density computation
e.g. Google TPU 3.0 pods are claimed to achieve 100PFLOPS
- On edge: low precision
e.g. int8 in cDSP supported by Qualcomm's SNPE
- Automatic kernel tuning & generation
An active research area. Typical projects include Halide³, TVM⁴ and TensorComprehension⁵


³<http://halide-lang.org/>

⁴<http://tvm-lang.org/>

⁵<https://facebookresearch.github.io/TensorComprehensions/>

- 1 Dense Numerical Computation
 - CPU Computation
 - Other Computation Devices
 - Computation & Memory Gap

Unbalanced Development of Processor and Memory



Graph from⁶

⁶Carlos Carvalho. "The gap between processor and memory speeds". In: *Proc. of IEEE International Conference on Control and Automation*. 2002.

Challenges from NN Architecture

Computation-sparse structure seems to be beneficial.

Architecture	Computation	Memory
Small kernel	$\frac{k^2}{k_1^2}$	Param $\frac{k_1^2}{k^2}$
Large stride	$\frac{1}{s^2}$	Output $\frac{1}{s^2}$
Group/depthwise conv	$\frac{1}{g^2}$	Param $\frac{1}{g}$
Shuffle/concat	0	1

Roofline Model

A visualization method to characterize computation/memory

Performance P (FLOPS) is approximately a function of arithmetic intensity I (FLOP/byte) for a particular architecture⁷.

Naïve Roofline

$$P = \min \left\{ \begin{array}{l} \pi \\ \beta \times I \end{array} \right.$$

where π is the peak computing performance and β is the peak bandwidth.

⁷Samuel Webb Williams. *Auto-tuning performance on multicore computers*.
University of California, Berkeley, 2008.

Roofline Model

A visualization method to characterize computation/memory




Image from https://en.wikipedia.org/wiki/Roofline_model

Thanks!

Questions and feedback are welcome :)