

Introduction to Computation Technologies in Deep Learning

Kai Jia
jiakai@megvii.com

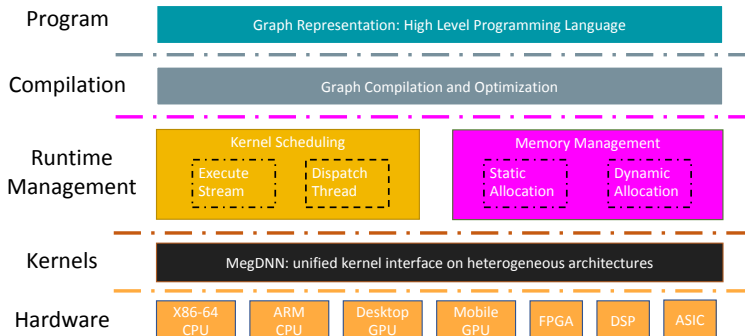
Megvii Inc.

March, 2019

- 1 Symbolic Computation
 - Representation
 - Execution & Optimization
- 2 Dense Numerical Computation
 - CPU Computation
 - Other Computation Devices
 - Computation & Memory Gap
- 3 Distributed Computation
 - System
 - Optimzation Algorithms
 - Communication Algorithms

Overview of a Deep Learning Framework

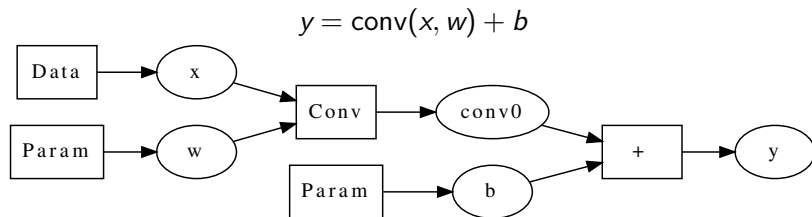
MegBrain Architecture



Computation Graph

$$y = \text{conv}(x, w) + b$$

Computation Graph



Graph Structure

Variable

- Corresponding to a tensor with concrete numerical values during graph execution

Graph Structure

Variable

- Corresponding to a tensor with concrete numerical values during graph execution
- Carrying some *attributes*:

DType Data type, like `int8` and `float32`.

Graph Structure

Variable

- Corresponding to a tensor with concrete numerical values during graph execution

- Carrying some *attributes*:

DType Data type, like `int8` and `float32`.

Shape Like `(5, 3)` for FC weight, `(128, 512, 7, 7)` for feature maps.

Example

Variable shape inference facilitates automatic weight initialization:

```
assert x.shape == (128, 50)
y = fully_connected(x, output_dim=100)
assert y.shape == (128, 100)
```

The weight matrix of this `FullyConnected` operator can be initialized to `np.random.normal((50, 100))`.

Graph Structure

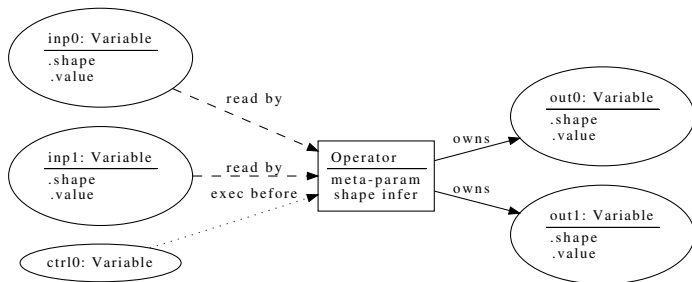
Operators & Edges

Operator

- Operators & variables form a bipartite graph
- Operators define the computation to be applied on input variables

Edge

- Data dependency: read input data
- Control dependency: require input operator to have finished



Operator Granularity

Trade-off between flexibility and ease-to-optimize

Category	Example	Advantage	Framework
Coarse-grained	$y = \text{BatchNorm}(x)$	Parsimony; Easy performance tuning	Caffe
Fine-grained	$y = \frac{x - \text{mean}(x)}{\text{std}(x)}$	Flexibility	Theano

Operator Granularity

Trade-off between flexibility and ease-to-optimize

Category	Example	Advantage	Framework
Coarse-grained	$y = \text{BatchNorm}(x)$	Parsimony; Easy performance tuning	Caffe
Fine-grained	$y = \frac{x - \text{mean}(x)}{\text{std}(x)}$	Flexibility	Theano

Our philosophy:

- Prefer **flexibility**: this can not be changed once the framework has been designed
- Utilize **multi-level API** for simplifying graph representation
- Speed can be continuously improved by **graph optimizer**

Auto Differentiation

Calculus 101: The Chain Rule

Straight-forward approach: each operator provides two methods: fprop and bprop.

```
inline Dtype Forward(const vector<Blob<Dtype>*>& bottom,  
    const vector<Blob<Dtype>*>& top);
```

```
inline void Backward(const vector<Blob<Dtype>*>& top,  
    const vector<bool>& propagate_down,  
    const vector<Blob<Dtype>*>& bottom);
```

Auto Differentiation

Calculus 101: The Chain Rule

Straight-forward approach: each operator provides two methods: fprop and bprop.

Limitations

- Graph optimizer can not be uniformly applied on both forward and backward passes
- Difficult to implement gradient of gradient ($\frac{\partial f(\frac{\partial L}{\partial x})}{\partial y}$, in WGAN training ¹) or higher-order gradients ($\frac{\partial^2 L}{\partial x^2}$).
- Difficult to modify/manipulate gradients (e.g. for low-bit training).

¹Ishaan Gulrajani et al. "Improved training of wasserstein gans". In: *arXiv preprint arXiv:1704.00028* (2017).

Auto Differentiation

Calculus 101: The Chain Rule

Unified approach: extending the graph with operators computing gradients of specific variables, via the chain rule.

$$y_1, \dots, y_m = f(x_1, \dots, x_n)$$

Gradient operator g for f .

Auto Differentiation

Calculus 101: The Chain Rule

Unified approach: extending the graph with operators computing gradients of specific variables, via the chain rule.

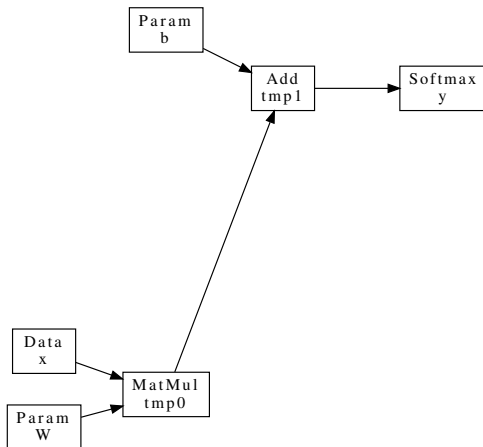
$$y_1, \dots, y_m = f(x_1, \dots, x_n)$$

Gradient operator g for f :

$$\frac{\partial L}{\partial x_1}, \dots, \frac{\partial L}{\partial x_n} = g\left(\frac{\partial L}{\partial y_1}, \dots, \frac{\partial L}{\partial y_m}, x_1, \dots, x_n, y_1, \dots, y_m\right)$$

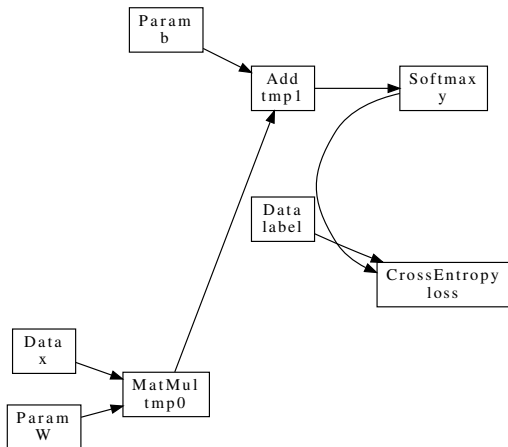
Auto Differentiation

Calculus 101: The Chain Rule



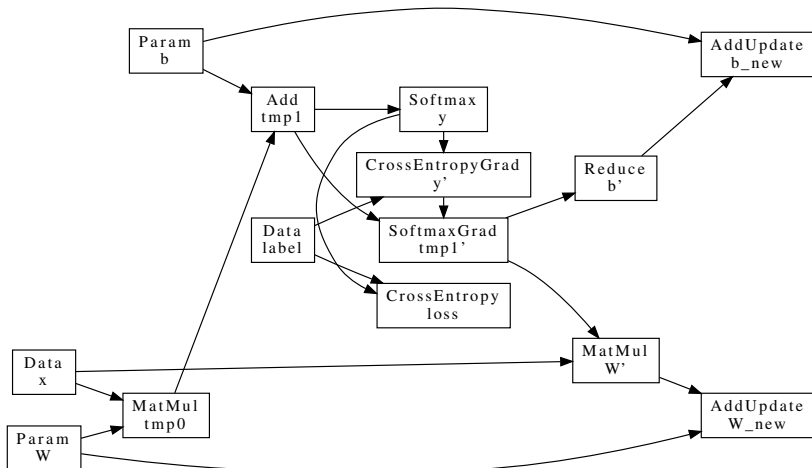
Auto Differentiation

Calculus 101: The Chain Rule



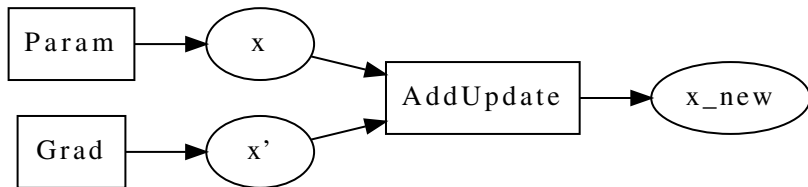
Auto Differentiation

Calculus 101: The Chain Rule



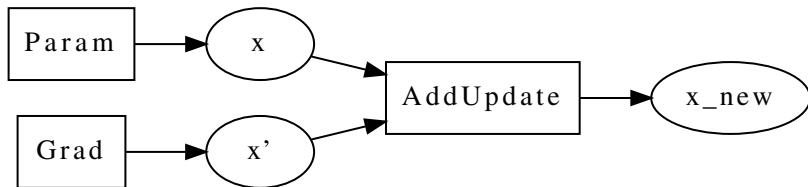
Mutable State

Enabling param updates to be expressed in graphs



Mutable State

Enabling param updates to be expressed in graphs

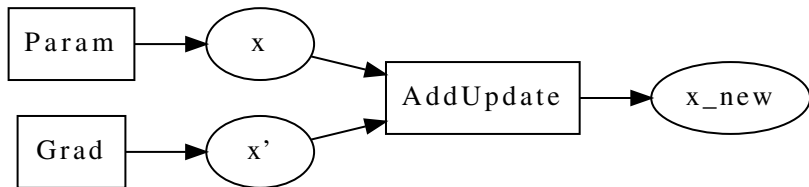


Note

- x and x_{new} share the underlying storage and should not be simultaneously read by one operator. Equivalently speaking, `AddUpdate` separates the graph.

Mutable State

Enabling param updates to be expressed in graphs



Note

- x and x_{new} share the underlying storage and should not be simultaneously read by one operator. Equivalently speaking, `AddUpdate` separates the graph.
- Readers of x must have finished (impl. by control dependency)

Symbolic Shape

Enabling computation involving tensor shapes

Example

Maxout¹:

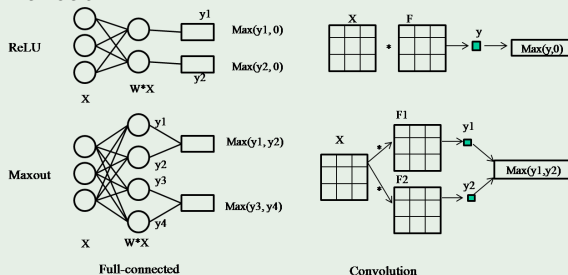


Image from ²

²Hai Dai Nguyen, Anh Duc Le, and Masaki Nakagawa. "Recognition of Online Handwritten Math Symbols Using Deep Neural Networks". In: *IEICE Trans. Inf. & Syst.* 99.12 (2016), pp. 3110–3118.

¹Ian J Goodfellow et al. "Maxout networks". In: *arXiv preprint*

Symbolic Shape

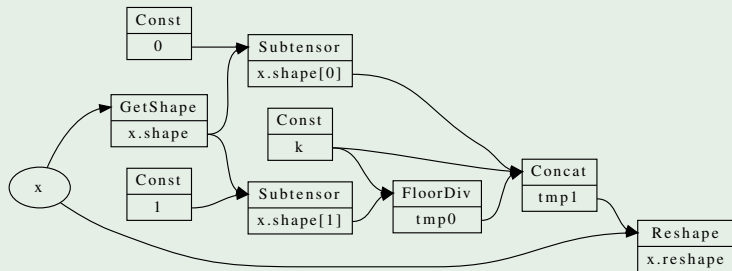
Enabling computation involving tensor shapes

Example

Maxout¹:

```
y = x.reshape(x.shape[0], x.shape[1] // k, k).max(axis=2)
```

where `x.shape` is also a symbol whose value is evaluated at runtime, so the computation adapts to different input shapes.



¹ Ian J Goodfellow et al. "Maxout networks". In: *arXiv preprint*

Symbolic Shape

Enabling computation involving tensor shapes

Example

Maxout¹:

```
y = x.reshape(x.shape[0], x.shape[1] // k, k).max(axis=2)
```

- Helps dealing with non-constant batch size or input image size
- Requires dynamic shape support: some shapes may remain unknown until graph execution

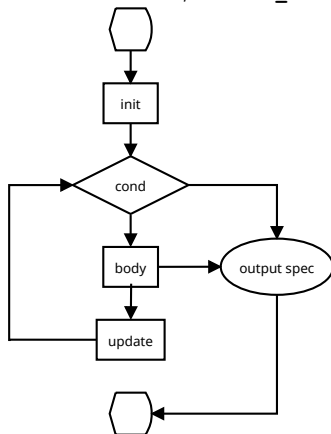
¹Ian J Goodfellow et al. "Maxout networks". In: *arXiv preprint arXiv:1302.4389* (2013).

Control Flow Operators

Towards universal computation (in theory)

Loop operator:

`scan` in Theano, `while_loop` in TensorFlow and `loop` in MegBrain.

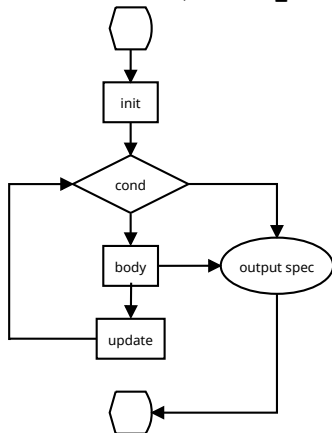


Control Flow Operators

Towards universal computation (in theory)

Loop operator:

`scan` in Theano, `while_loop` in TensorFlow and `loop` in MegBrain.



- With symbolic shapes and control flow operators, a computation graph is **Turing-complete**!
- Useful for RNN and iterative algorithms

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Static Computation Graph

- **Unfamiliar programming model:**

Stateless, functional: `y = x.setsub[1:3](xs)`

rather than imperative: `x[1:3] = xs`

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Static Computation Graph

- **Unfamiliar programming model:**

Stateless, functional: `y = x.setsub[1:3](xs)`

rather than imperative: `x[1:3] = xs`

- **Difficult to debug:** code is written for graph construction but tensor values can only be known during graph execution

`y = printop(y)` rather than `print(y)`

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

- Auto differentiation: keep symbolic track of computation path

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

- Auto differentiation: keep symbolic track of computation path
- Difficulties:
 - Hard to optimize: lack of global information
 - Hard to deploy: graph depends on code

Dynamic Computation Graph

Ease of programming beyond Turing-completeness

Dynamic Computation Graph

- Implemented by eager evaluation:

```
while (a.dot(x) - I).max().getvalue() > eps:  
    x = x.dot(2 * I - a.dot(x))  
print(grad(loss, x).getvalue())
```

- Auto differentiation: keep symbolic track of computation path
- Difficulties:
 - Hard to optimize: lack of global information
 - Hard to deploy: graph depends on code
 - Solution: tracing, source code transformation ²

²Dan Moldovan et al. "AutoGraph: Imperative-style Coding with Graph-based Performance". In: *SysML* (2019).

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Graph Execution

Separation of representation and execution allows abstraction of hardware details

- Map from variables to tensor values

Graph Execution

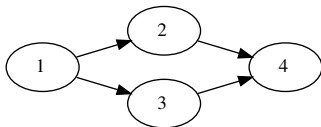
Separation of representation and execution allows abstraction of hardware details

- Map from variables to tensor values
- Map from operators to kernels on some specific architecture

Graph Execution

Separation of representation and execution allows abstraction of hardware details

- Map from variables to tensor values
- Map from operators to kernels on some specific architecture
- Kernels are scheduled according to topological order



(1, 2, 3, 4) or (1, 3, 2, 4)

Optimizing by Graph Transformation

- Expression simplifying: $x + 1 - 2 + x \Rightarrow 2x - 1$

Optimizing by Graph Transformation

- Expression simplifying: $x + 1 - 2 + x \Rightarrow 2x - 1$
- Operation reordering according to shapes
tensors: x and y
scalars: a and b
 $x + a + y + b \Rightarrow a + b + x + y$

Optimizing by Graph Transformation

- Expression simplifying: $x + 1 - 2 + x \Rightarrow 2x - 1$
- Operation reordering according to shapes
tensors: x and y
scalars: a and b
 $x + a + y + b \Rightarrow a + b + x + y$
- Operator fusion: $x \cdot y + z \Rightarrow \text{fma}(x, y, z)$
 - static fusion: predefined fusion rules
 - dynamic fusion: Just-in-time compilation (JIT) for actual computation graph

Runtime Memory Management

- Baseline: reference counting + some classical memory allocator

Runtime Memory Management

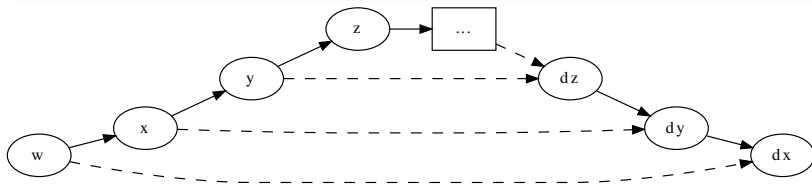
- Baseline: reference counting + some classical memory allocator
 - Readonly forwarding: reuse input storage for operators like reshape and subtensor
 - Writable forwarding (a.k.a. inplace operation): overwrite input storage
- Caution:** must ensure no other readers exist (i.e. refcnt equals 1)

Sublinear Memory

Trade time for memory

Observation

- Long-term dependency for gradient computing consumes lots of memory.
- Assume $x_{i+1} = \text{conv}(x_i, w_i)$, then x_{i+1} can only be discarded after $\frac{\partial L}{\partial w_i}$ and $\frac{\partial L}{\partial x_i}$ have been computed.

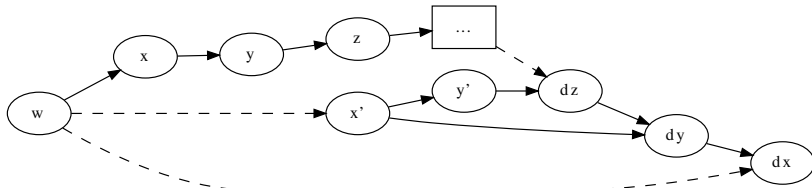


Sublinear Memory

Trade time for memory

Observation

- Long-term dependency for gradient computing consumes lots of memory.
- Assume $x_{i+1} = \text{conv}(x_i, w_i)$, then x_{i+1} can only be discarded after $\frac{\partial L}{\partial w_i}$ and $\frac{\partial L}{\partial x_i}$ have been computed.



Sublinear Memory

Trade time for memory

Observation

- Long-term dependency for gradient computing consumes lots of memory.
- Assume $x_{i+1} = \text{conv}(x_i, w_i)$, then x_{i+1} can only be discarded after $\frac{\partial L}{\partial w_i}$ and $\frac{\partial L}{\partial x_i}$ have been computed.

Method

- Split the sequence into blocks consisting of consecutive operators and keep only the first variable in any block
- Recompute internal values in a block when gradient is needed
- In the above example, discard x_{km+j} for all $0 < j < m$ and recompute them when needed.

Sublinear Memory

Trade time for memory

Reduce memory usage to $O(\sqrt{n})$ with extra $O(n)$ time cost in the ideal case.

For a graph with 10000 convolutions and their gradients:

comp_node	alloc	lower_bound	upper_bound
gpu0:0	15624.37MiB(16383336448bytes)	15624.37MiB(100.00%)	31889.13MiB(204.10%)
comp_node	alloc	lower_bound	upper_bound
gpu0:0	173.03MiB(181430784bytes)	168.76MiB(97.53%)	47251.78MiB(27309.08%)

Note: this idea is also published in³.

³Tianqi Chen et al. "Training deep nets with sublinear memory cost". In: *arXiv preprint arXiv:1604.06174* (2016).

- 1 Symbolic Computation
 - Representation
 - Execution & Optimization
- 2 Dense Numerical Computation
 - CPU Computation
 - Other Computation Devices
 - Computation & Memory Gap
- 3 Distributed Computation
 - System
 - Optimzation Algorithms
 - Communication Algorithms

Instruction: The Hardware/Software Interface

What is a program?

```
int sum(int *x) {  
    return x[0] + x[1];  
}
```

00000000000000000 <sum>:

0: 8b 47 04

3: 03 07

5: c3

mov 0x4(%rdi),%eax

add (%rdi),%eax

retq

Modern CPU Technologies

Instr. No.	Pipeline Stage						
	IF	ID	EX	MEM	WB		
1							
2							
3							
4							
5							
Clock Cycle	1	2	3	4	5	6	7

- **Pipeline** ⁴
- **Superpipelining** increases stage number and simplifies each stage
- **Superscalar** dispatches multiple instructions to implement instruction-level parallelism
- **Out-of-order execution** executes according to availability of input data rather than original program order

⁴ Image from https://en.wikipedia.org/wiki/Instruction_pipelining

SIMD

Single instruction, multiple data

Store multiple data items in one register and process them in a single instruction.

Calculation of theoretical FLOPS⁵

$$FLOPS = f \cdot w \cdot IPC$$

f : frequency

w : SIMD width (number of floats per register)

IPC : SIMD instructions per cycle

⁵floating point operations per second

SIMD

Single instruction, multiple data

Example

Intel® CPUs usually have $IPC = 2$. However if FMA is supported, IPC should be counted as 4 since 2 FMA instructions is essentially 4 floating point operations.

Intel® Xeon® Platinum 8180M⁵

# of Cores	28
Processor Base Frequency	2.50 GHz
Max Turbo Frequency	3.80 GHz
# of AVX-512 FMA Units	2

$$\begin{aligned}
 FLOPS &= 3.8 \text{Gcyc/s} \times 4 \text{instr/cyc} \times 16 \text{float/instr} \\
 &= 243.2 \text{GFLOPS}
 \end{aligned}$$

$$FLOPS_{TOT} = FLOPS \times 28 = 6.8 \text{TFLOPS}$$

⁵ data available at

Memory Hierarchy

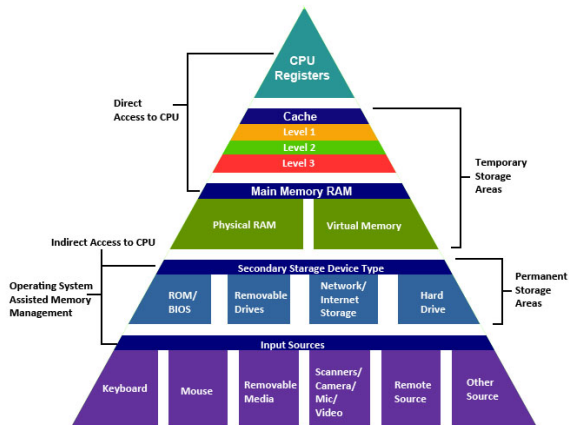


Image from <http://cse1.net/recaps/4-memory.html>

Static Random-Access Memory (SRAM)

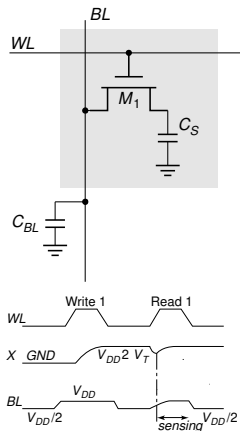


- 1 Fast
- 2 Low power consumption
- 3 No refresh circuit

Image from https://en.wikipedia.org/wiki/Static_random-access_memory

RAM Implementation

Dynamic Random-Access Memory (DRAM)



Advantages:

- ① High density
- ② Cheap

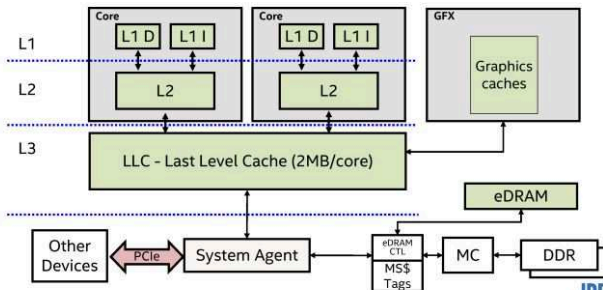
Refresh: periodically read blocks and write back.

Image from <http://docencia.ac.upc.edu/master/MIRI/NCD/docs/04-Memory%20Structures-2.pdf>

Cache Hierarchy

A hierarchical design for better trade-off between memory capacity and latency.

eDRAM Based Cache



26

Intel Next Generation Microarchitecture Code Name Skylake

IDF15
INTEL DEVELOPER FORUM

Image from

<https://www.anandtech.com/show/9582/intel-skylake-mobile-desktop-launch-architecture-analysis/5>

Cache Hierarchy

A hierarchical design for better trade-off between memory capacity and latency.

Core i7 Xeon 5500 Series

L1 hit	~ 4 cycles
L2 hit	~ 10 cycles
L3 hit line unshared	~ 40 cycles
L3 hit, shared line in another core	~ 65 cycles
L3 hit, modified in another core	~ 75 cycles
Remote L3	~ 100 – 300 cycles
Local DRAM	~ 60 ns
Remote DRAM	~ 100 ns

source:

https://software.intel.com/sites/products/collateral/hpc/vtune/performance_analysis_guide.pdf

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

- Associativity

Number of different tags to be kept under the same index

CPU Cache Structure

- Cache line

tag	data block	flag bits (valid, dirty)
-----	------------	--------------------------

- Indexing

tag (40bit)	index (6bit)	block offset (6bit)
-------------	--------------	---------------------

- Associativity

Number of different tags to be kept under the same index

- VIPT Addressing

Virtually indexed, physically tagged (VIPT)

- Solve aliasing problem
- Simultaneous cache and TLB lookup

Interesting reading: <http://igoro.com/archive/gallery-of-processor-cache-effects/>

CPU Cache Structure

Example

```
$ grep -m1 name /proc/cpuinfo
model name      : Intel(R) Core(TM) i5-6200U CPU @ 2.30GHz

$ cat /sys/devices/system/cpu/cpu0/cache/index0/{size,ways_of_associativity,coherency_line_size}
32K
8
64
```

- block offset: $\log_2 64 = 6\text{bit}$
- index: $\log_2(32\text{KiB}/64\text{B}/8) = 6\text{bit}$
- tag: $52 - 6 - 6 = 40\text{bit}$ (48-bit virtual memory and 52-bit physical memory)

Note that *block offset* and *index* together take 12 bits, which is equal to page size (4KiB), so VIPT can be easily implemented.

A MatMul Example

```
void matmul(float *a, float *b, float *c, int n) {  
    for (int i = 0; i < n; ++ i) {  
        for (int j = 0; j < n; ++ j) {  
            float sum = 0;  
            for (int k = 0; k < n; ++ k) {  
                sum += a[i * n + k] * b[k * n + j];  
            }  
            c[i * n + j] = sum;  
        }  
    }  
}
```

A MatMul Example

```
void matmul(float *a, float *b, float *c, int n) {  
    for (int i = 0; i < n; ++ i) {  
        for (int j = 0; j < n; ++ j) {  
            float sum = 0;  
            for (int k = 0; k < n; ++ k) {  
                sum += a[i * n + k] * b[k * n + j];  
            }  
            c[i * n + j] = sum;  
        }  
    }  
}
```

Swap the loops on j and k

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

NVIDIA GPU

A single instruction, multiple thread architecture

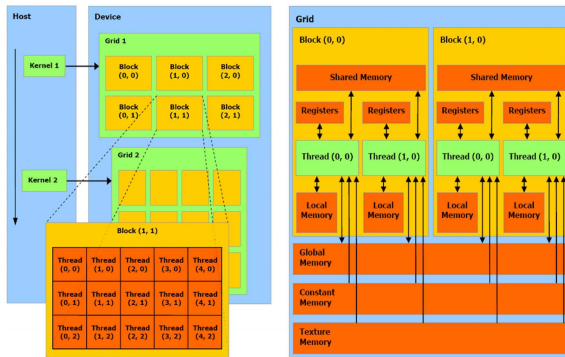


Image from⁶

⁶Marco Nobile et al. "cuTauLeaping: A GPU-Powered Tau-Leaping Stochastic Simulator for Massive Parallel Analyses of Biological Systems". In: 9 (Mar. 2014), e91963.

NVIDIA GPU

A single instruction, multiple thread architecture

```
__global__ void add(float *a, float *b, float *c, int n) {  
    int id = blockIdx.x*blockDim.x+threadIdx.x;  
    if (id < n)  
        c[id] = a[id] + b[id];  
}
```

- **Memory Coalescing**

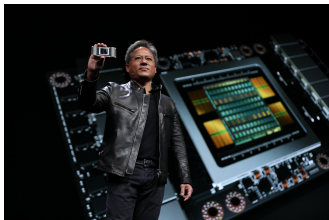
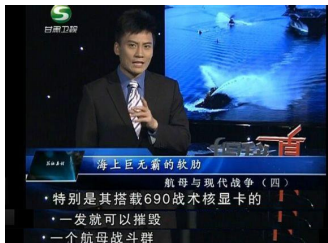
Adjacent threads access adjacent memories simultaneously

- **Parallelism**

Divide the total work among many tiny threads

NVIDIA GPU

A single instruction, multiple thread architecture



Tesla V100 for NVLink @ 300W

- 15.7 *TFLOPS* for single-precision
- **125 *TFLOPS*** for half-precision

Tesla T4 @ 75W

- 130T INT8
- 260T INT4

The Trend

- On cloud: high density computation
e.g. Google TPU 3.0 pods are claimed to achieve 100PFLOPS
- On edge: low precision
e.g. int8 in cDSP supported by Qualcomm's SNPE
- Automatic kernel tuning & generation
An active research area. Typical projects include Halide⁶, TVM⁷ and TensorComprehension⁸

⁶<http://halide-lang.org/>

⁷<http://tvm-lang.org/>

⁸<https://facebookresearch.github.io/TensorComprehensions/>

1 Symbolic Computation

- Representation
- Execution & Optimization

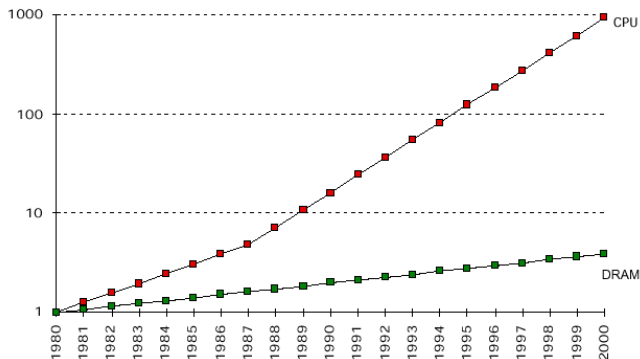
2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Unbalanced Development of Processor and Memory



Graph from⁹

⁹Carlos Carvalho. "The gap between processor and memory speeds". In: *Proc. of IEEE International Conference on Control and Automation*. 2002.

Challenges from NN Architecture

Computation-sparse structure seems to be beneficial.

Architecture	Computation	Memory
Small kernel	$\frac{k^2}{k_1^2}$	Param $\frac{k_1^2}{k^2}$
Large stride	$\frac{1}{s^2}$	Output $\frac{1}{s^2}$
Group/depthwise conv	$\frac{1}{g^2}$	Param $\frac{1}{g}$
Shuffle/concat	0	1

Roofline Model

A visualization method to characterize computation/memory

Performance P (FLOPS) is approximately a function of arithmetic intensity I (FLOP/byte) for a particular architecture¹⁰.

Naïve Roofline

$$P = \min \left\{ \begin{array}{l} \pi \\ \beta \times I \end{array} \right.$$

where π is the peak computing performance and β is the peak bandwidth.

¹⁰Samuel Webb Williams. *Auto-tuning performance on multicore computers*.
University of California, Berkeley, 2008.

Roofline Model

A visualization method to characterize computation/memory

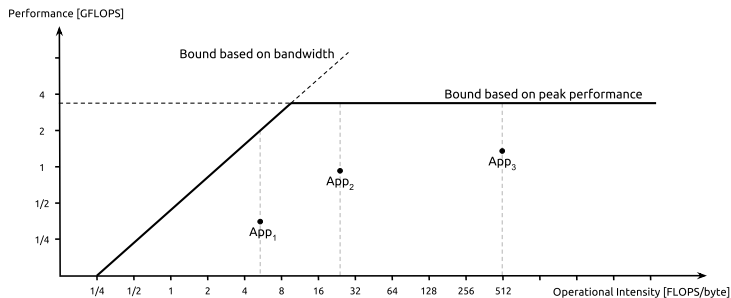


Image from https://en.wikipedia.org/wiki/Roofline_model

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

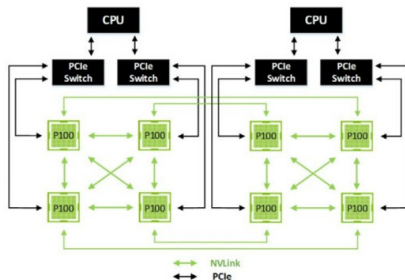
- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Communication System: Single Node

- PCI-e: connection between GPUs, network adaptors and others
 - Switches may be needed
 - 985 MiB/s each PCI-e 3.0 lane
 - LGA-2011 socket: 40 lanes
- NVLink: GPU interconnect by NVIDIA



Communication System: LAN

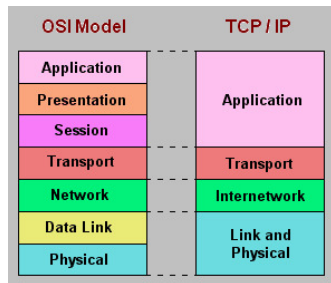


Image from <http://www.just2good.co.uk/tcpipStack.php>

- Ethernet: 10M to 100G, latency 100 - 20 μ s
- InfiniBand: 2.5 to 250G, latency 5 - 0.5 μ s

RDMA

Remote Direct Memory Access

Bypass the TCP/IP stack and free CPU from handling packets.

RDMA

Remote Direct Memory Access

Bypass the TCP/IP stack and free CPU from handling packets.

- RoCE: RDMA over Converged Ethernet
- InfiniBand: RDMA supported
- NVIDIA GPUDirect: RDMA between GPUs

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

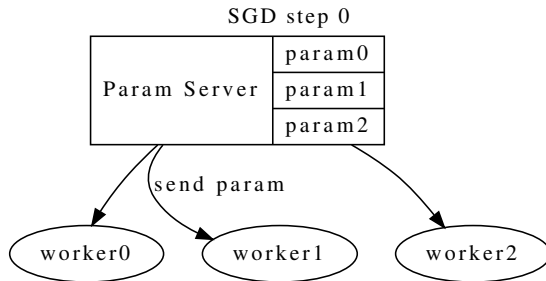
$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$

Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$

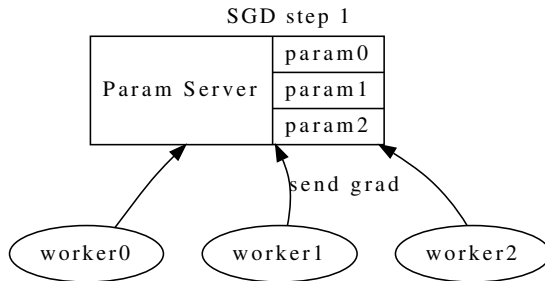


Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$



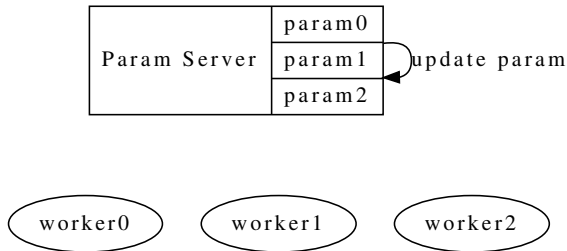
Synchronous SGD

Basics

Each worker processes a part of the whole batch; params and grads must be synchronized

$$L(\{d_0, d_1\}, W) = \alpha_0 L(\{d_0\}, W) + \alpha_1 L(\{d_1\}, W)$$

SGD step 2



Asynchronous SGD

Basics

Each worker has an outdated local copy of params and updates central param storage asynchronously. Friendly for parallel speedup.

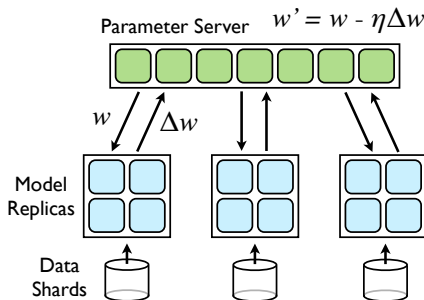


Image from¹⁰

¹⁰Jeffrey Dean et al. "Large scale distributed deep networks". In: *NIPS*. 2012, pp. 1223–1231.

Asynchronous SGD

Difficulties

ASGD is not equivalent to SGD and it is hard to tune due to noisy gradients. Many works exist on analyzing convergence and improving performance¹¹¹²¹³.

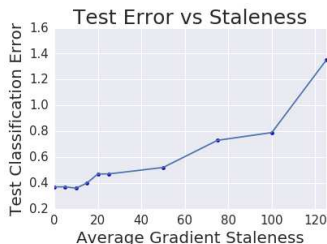


Image from¹⁴

¹¹Xiangru Lian et al. "Asynchronous Parallel Stochastic Gradient for Nonconvex Optimization". In: *NIPS*. 2015, pp. 2737–2745.

¹²Wei Zhang et al. "Staleness-aware async-SGD for Distributed Deep Learning". In: *IJCAI*. 2016, pp. 2350–2356.

¹³Sixin Zhang, Anna E Choromanska, and Yann LeCun. "Deep learning with elastic averaging SGD". In: *NIPS*. 2015, pp. 685–693.

¹⁴Jianmin Chen et al. "Revisiting distributed synchronous SGD". In: *arXiv preprint arXiv:1604.00981* (2016).

Synchronous SGD

Reduce communication by compressing gradients¹⁵

$$Q_s(v_i) = \|\mathbf{v}\|_2 \cdot \text{sgn}(v_i) \cdot \xi_i(\mathbf{v}, s)$$

¹⁵Ryota Tomioka and Milan Vojnovic. "QSGD: Communication-Efficient Stochastic Gradient Descent, with Applications to Training Neural Networks".

In: *arXiv preprint arXiv:1610.02132* (2016).

Synchronous SGD

Handle straggling workers by backup workers¹⁶

Use $N + b$ workers but only receive gradients from any N of them and do not wait for the slowest b workers.

¹⁶Jianmin Chen et al. "Revisiting distributed synchronous SGD". In: *arXiv preprint arXiv:1604.00981* (2016).

Synchronous SGD

Ensure accuracy by careful hyperparam tuning

8192 minibatch size on 256 GPUs¹⁷

$$\hat{\eta} = k\eta$$
$$m = \frac{\eta_{t+1}}{\eta_t}$$

32K minibatch size with LARS¹⁸

$$\lambda^l = \eta \times \frac{\|w^l\|}{\|\Delta L(w^l)\|}$$

¹⁷Priya Goyal et al. "Accurate, Large Minibatch SGD: Training ImageNet in 1 Hour". In: *arXiv preprint arXiv:1706.02677* (2017).

¹⁸Yang You, Igor Gitman, and Boris Ginsburg. "Large batch training of convolutional networks". In: *arXiv preprint arXiv:1708.03888* (2017).

1 Symbolic Computation

- Representation
- Execution & Optimization

2 Dense Numerical Computation

- CPU Computation
- Other Computation Devices
- Computation & Memory Gap

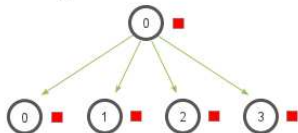
3 Distributed Computation

- System
- Optimzation Algorithms
- Communication Algorithms

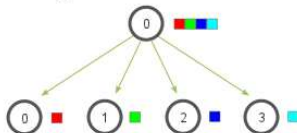
MPI Primitives

Collective communication routines in MPI are common in distributed DL

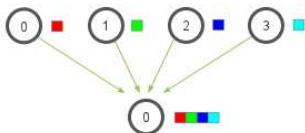
MPI_Bcast



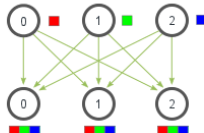
MPI_Scatter



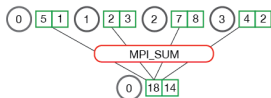
MPI_Gather



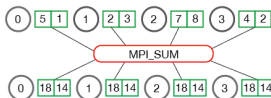
MPI_Allgather



MPI_Reduce



MPI_Allreduce



An AllReduce Algorithm

Assume message size K and number of workers N

- Reduce to a worker (assume W_{N-1} here):
 W_i sends to W_{i+1} at step i ; communication per worker is N
- Broadcast from a worker: as above

An AllReduce Algorithm


Assume message size K and number of workers N

- Reduce to a worker (assume W_{N-1} here):
 W_i sends to W_{i+1} at step i ; communication per worker is N
- Broadcast from a worker: as above
- AllReduce:
 - 1 Split the message into N parts
 - 2 Reduce the i th part to W_i ; all reductions run in parallel
 - 3 Broadcast each reduced part to all workers in parallel

Communication cost for each worker is $2(N-1)\frac{K}{N}$,
independent of N .

More details and discussions are given in^{19 20}.

¹⁹Rajeev Thakur, Rolf Rabenseifner, and William Gropp. "Optimization of collective communication operations in MPICH". In: *The International Journal of High Performance Computing Applications* 19.1 (2005), pp. 49–66.

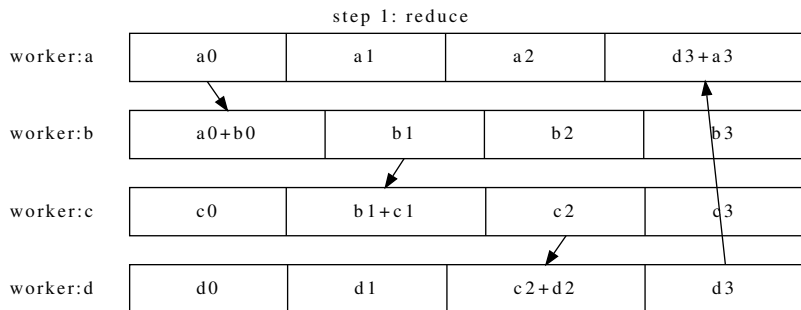
²⁰<http://research.baidu.com/bringing-hpc-techniques-deep-learning> 

An AllReduce Algorithm

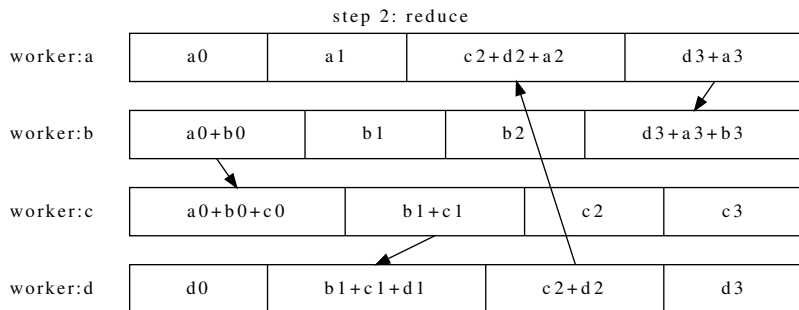
step 0: init

worker:a	a0	a1	a2	a3
worker:b	b0	b1	b2	b3
worker:c	c0	c1	c2	c3
worker:d	d0	d1	d2	d3

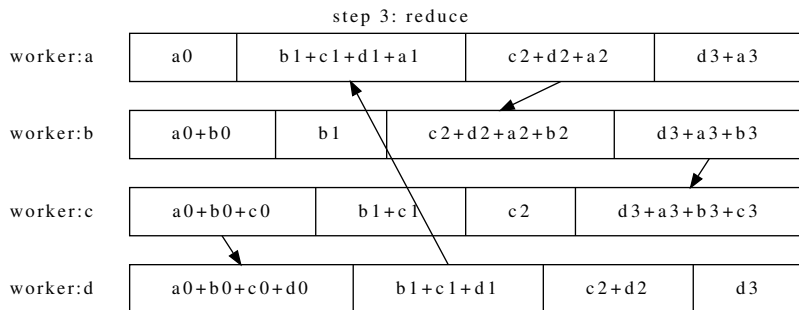
An AllReduce Algorithm



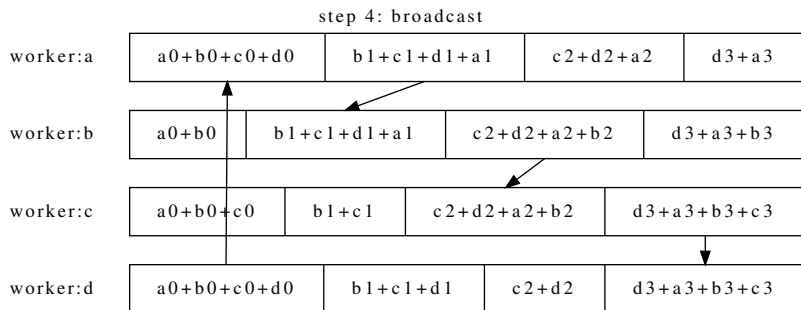
An AllReduce Algorithm



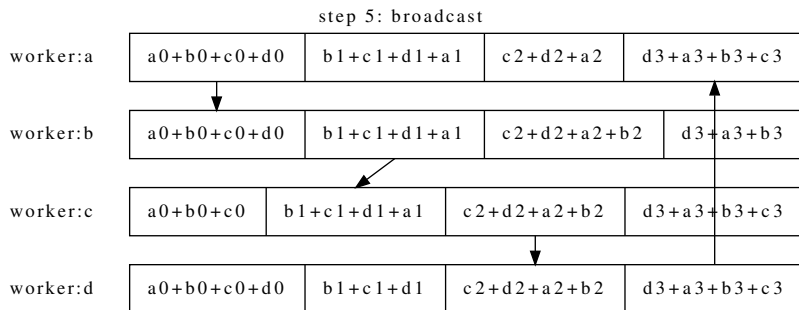
An AllReduce Algorithm



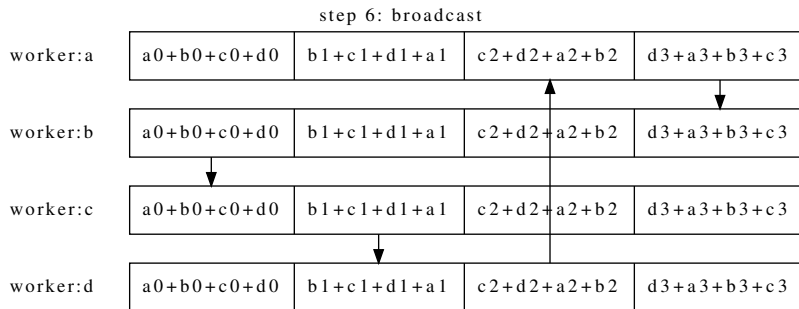
An AllReduce Algorithm



An AllReduce Algorithm



An AllReduce Algorithm



Thanks!

Questions and feedback are welcome :)